

Principal Component Analysis

CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez
University of Rhode Island

Dimensionality reduction

- Basically ...
 - ✓ mapping high-dimensional data into low dimensional data
 - ✓ can be as simple as dropping some dimensions or using a combination of all dimensions
- Given $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ with $\mathbf{x}_i \in \mathbb{R}^d$, find a representation $\mathcal{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_n\}$ with $\mathbf{z}_i \in \mathbb{R}^{d'}$ and $d' \ll d$
 - ✓ properties should be preserved (e.g., variance, distances, neighborhood)

$$\begin{bmatrix} \text{---} & \mathbf{x}_1^T & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_n^T & \text{---} \end{bmatrix}_{n \times d} \Rightarrow \begin{bmatrix} \mathbf{z}_n^T \\ \vdots \\ \mathbf{z}_n^T \end{bmatrix}_{n \times d'}$$


Why dimensionality reduction?

- Data visualization (2-d or 3-d)
- Preprocess data before machine learning
 - ✓ algorithms can focus on important features/patterns
 - ✓ training can be more efficient
- Removing noise and redundant information
- Data compression

Principal component analysis (eigendecomposition formulation)


Principal component analysis

- Find projections of the data onto directions that maximize variance
- ✓ directions are orthogonal to each other



<https://www.dataschool.io/15-hours-of-expert-machine-learning-videos/>

Preliminaries



The **vector projection** of vector \mathbf{a} onto a nonzero vector \mathbf{b} is the orthogonal projection of \mathbf{a} onto a straight line parallel to \mathbf{b} .


$$\mathbf{a}_1 = a_1 \hat{\mathbf{b}}$$

where a_1 is a scalar, called the **scalar projection** of \mathbf{a} onto \mathbf{b} , and $\hat{\mathbf{b}}$ is the unit vector in the direction of \mathbf{b} . The scalar projection is defined as:

$$a_1 = \mathbf{a} \cdot \hat{\mathbf{b}}$$


https://en.wikipedia.org/wiki/Vector_projection

Data projection on a line



Idea: finding a line (subspace) that maintains as much variance (spread) as possible when the data is projected onto this subspace.

What is the best line?



Idea: minimize sum of square distances to the line (or maximize the sum of squares of the scalar projections)

Disclaimer


► The following slides contain figures adapted from:

✓ Unsupervised Learning **SERRANO.ACADEMY**
the art of understanding
✓ <https://serrano.academy/unsupervised-learning/>


► Video:

✓ <https://www.youtube.com/watch?v=g-Hb26agBFg>


The mean




Doesn't always help



The variance




$$\text{Variance} = \frac{1^2 + 0^2 + 1^2}{3} = 2/3$$




$$\text{Variance} = \frac{5^2 + 0^2 + 5^2}{3} = 10/3$$

Variance and data




Doesn't always help




$$\text{x-variance} = \frac{2^2 + 0^2 + 2^2}{3} = 8/3$$

$$\text{y-variance} = \frac{1^2 + 0^2 + 1^2}{3} = 2/3$$

Covariance



$$\text{covariance} = \frac{(-2) + 0 + (-2)}{3} = -4/3$$



$$\text{covariance} = \frac{2 + 0 + 2}{3} = 4/3$$

Covariance



negative
covariance




covariance zero
(or very small)




positive
covariance


Covariance matrix



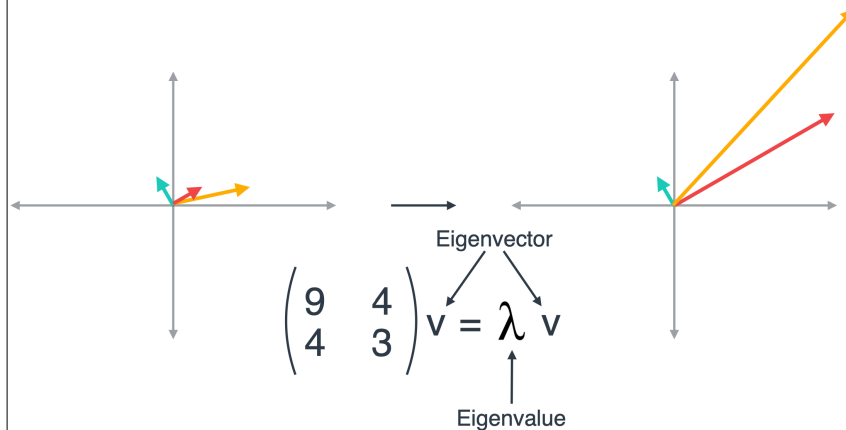
Linear transformations




Eigenvectors and eigenvalues




Eigenvectors and eigenvalues




Principal component analysis (PCA)




Principal component analysis (PCA)



Principal component analysis (PCA)



Principal component analysis (PCA)



More formally

- Given $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ with $\mathbf{x}_i \in \mathbb{R}^d$
- Find $\mathcal{Z} = \{\mathbf{z}_1, \dots, \mathbf{z}_n\}$ with $\mathbf{z}_i \in \mathbb{R}^{d'}$ and $d' \ll d$
- Example from 3-d to 2-d

$$\mathbf{u}_1 \in \mathbb{R}^3, \mathbf{u}_2 \in \mathbb{R}^3$$

$$\mathbf{z}_i = [\mathbf{u}_1^T \mathbf{x}_i, \mathbf{u}_2^T \mathbf{x}_i], \mathbf{z}_i \in \mathbb{R}^2$$

PCA

▸ Input

✓ data must be centered

$$X = \begin{bmatrix} \text{---} & \mathbf{x}_1^T & \text{---} \\ & \vdots & \\ \text{---} & \mathbf{x}_n^T & \text{---} \end{bmatrix}_{n \times d}$$

$$\frac{1}{n} \sum_{i=1}^n c_i = 0, \quad \text{for all columns } c \text{ in } X$$

▸ Output

✓ orthonormal principal components $\mathbf{u}_1, \dots, \mathbf{u}_k$

Eigendecomposition of symmetric matrices

As a special case, for every $n \times n$ real symmetric matrix, the eigenvalues are real and the eigenvectors can be chosen real and orthonormal. Thus a real symmetric matrix A can be decomposed as:

$$A = Q\Lambda Q^T$$

where Q is a matrix whose columns are eigenvectors of A , and Λ is a diagonal matrix whose entries are the eigenvalues of A .

https://en.wikipedia.org/wiki/Eigendecomposition_of_a_matrix

First principal component

$$\|\mathbf{u}_1\|_2 = 1$$

$$\arg \max_{\mathbf{u}_1} \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{u}_1)^2$$

$$\arg \max_{\mathbf{u}_1} \mathbf{u}_1^T \left[\sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^T \right] \mathbf{u}_1$$

XX^T

solution \mathbf{u}_1 is the first eigenvector of XX^T

Second principal component

$$\|\mathbf{u}_2\|_2 = 1, \quad \mathbf{u}_2^T \mathbf{u}_1 = 0$$

$$\arg \max_{\mathbf{u}_2} \sum_{i=1}^n (\mathbf{x}_i^T \mathbf{u}_2)^2$$

$$\arg \max_{\mathbf{u}_2} \mathbf{u}_2^T (XX^T) \mathbf{u}_2$$

solution \mathbf{u}_2 is the second eigenvector of XX^T

PCA algorithm

- ▶ Assuming X and K as inputs
 - ✓ Center the data in X
 - ✓ Compute eigendecomposition $(U\Lambda U^T)$ of XX^T
 - ✓ Return the top K eigenvectors from U
- ▶ Eigenvectors can be use for projecting the data in lower dimensions