

# Linear Regression


CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez  
University of Rhode Island

## Continuous targets

- Certain applications require the prediction of continuous values





<https://towardsdatascience.com/stock-prediction-in-python-b66555171a2>

## Linear model

- Assumes the output  $y$  is a linear function of the input  $x$

✓ can use the function to make predictions, very simple approach, e.g. linear regression




$$y = wx + b \quad w, x, b \in \mathbb{R}$$

## Linear model

- What if we have  $d$  features?

why transpose?



$$y = \mathbf{w}^T \mathbf{x} + b \quad \mathbf{w}, \mathbf{x} \in \mathbb{R}^d \quad b \in \mathbb{R}$$

Figure from <https://www.statlearning.com/>

## Linear model

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

hypothesis


weights

bias


The weights and bias are the model **parameters** which define the hypothesis and are used to make predictions

## Draw the models


$$y = w_0 + w_1 x_1 \quad w_0 = b$$



$$\mathbf{w} = [0.5, 0]$$




$$\mathbf{w} = [0, 1.5]$$



$$\mathbf{w} = [0.5, 0.5]$$


## Residuals




Want a **linear function** with **small residuals**

<http://work.caltech.edu/slides/slides03.pdf>

## Residuals and Loss



High loss



Low loss

From <https://developers.google.com/machine-learning/crash-course/descending-into-ml/training-and-loss>

## Goal of learning

- Find a function that best approximates target function (**minimize expected loss**)

For  $h \in \mathcal{H}$  and  $\forall (\mathbf{x}^{(i)}, y^{(i)}) \sim P$ , we want  $h(\mathbf{x}^{(i)}) \approx f(\mathbf{x}^{(i)})$

- What is the **expected loss**?

✓ cannot calculate, can **approximate** with empirical loss:

$$\mathbb{E} [l(h, \mathbf{x}^{(i)}, y^{(i)})]_{(\mathbf{x}^{(i)}, y^{(i)}) \sim P} \approx L(h, \mathcal{D})$$

## Defining linear regression

- Data  $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$   
 $\mathbf{x}^{(i)} \in \mathbb{R}^d, y^{(i)} \in \mathbb{R}$

- Loss Function: Squared Loss (MSE)

$$l_{sq}(h, \mathbf{x}, y) = (h(\mathbf{x}) - y)^2$$

$$L_{sq}(h, \mathcal{D}) = \frac{1}{n} \sum_{i=1}^n l_{sq}(h, \mathbf{x}^{(i)}, y^{(i)})$$

## Alternative notation

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

$$= \sum_{i=1}^d w_i x_i + b$$

$$= \sum_{i=0}^d w_i x_i \quad x_0 = 1, \quad w_0 = -b$$

$$= \mathbf{w}^T \mathbf{x}$$

x0	x1	x2	y
1	0.5	0.1	0.25
1	0.3	0.9	0.5
1	0.3	0.875	1.15
1	0.45	0.15	2.13
...	...	...	...

## Alternative notation

input:  $(x_0, x_1, \dots, x_d)$

model:  $(w_0, w_1, \dots, w_d)$

$$h(\mathbf{x}) = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \mathbf{x}$$


What is the hypothesis space?

all linear functions in  $\mathbb{R}^{d+1}$

$$\mathcal{H} = \{h_w : h_w(\mathbf{x}) = \mathbf{w}^T \mathbf{x}, \mathbf{w} \in \mathbb{R}^{d+1}\}$$

# Closed-form solution

## min vs. argmin



$\min f(x)?$

$\arg \min_x f(x)?$

$$f(x) = 6x^2 + 4x + 5$$

## Solving linear regression (least squares)

find these parameters

$$\mathbf{w}^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$$

given this objective function

unconstrained optimization

## Solving linear regression (least squares)

$$\mathbf{w}^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$$

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$

## Understanding matrix form

$$\mathbf{y} = \mathbf{X}\mathbf{w}$$

$$\begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \approx \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ \vdots \\ (\mathbf{x}^{(n)})^T \end{bmatrix} \cdot \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix}$$

input vector  $\mathbf{x}^{(1)}$

## Norms

- ▶ A **norm** is a function that assigns a strictly positive length to each vector in a vector space
  - ✓ except for the zero vector

$\ell_1$ -norm:  
manhattan distance  
from origin

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|$$

$\ell_2$ -norm:  
euclidean norm, euclidean  
distance from origin

$$\|\mathbf{x}\|_2 = \left( \sum_{i=1}^n x_i^2 \right)^{\frac{1}{2}}$$


## Using matrix notation

$$\arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$$

$$\arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

## How to minimize it?


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$



continuous,  
differentiable,  
**convex**

**optimal  
solution**

## (some) Critical points



[https://en.wikipedia.org/wiki/Maxima\\_and\\_minima](https://en.wikipedia.org/wiki/Maxima_and_minima)

## Closed form solution

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E_{\text{in}}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} \quad \text{where} \quad \mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

$\mathbf{X}^\dagger$  is the 'pseudo-inverse' of  $\mathbf{X}$

There are other methods for finding the optimal solution  
e.g. gradient descent, MLE

<http://work.caltech.edu/slides/slides03.pdf>

## The algorithm

- 1: Construct the matrix  $\mathbf{X}$  and the vector  $\mathbf{y}$  from the data set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$  as follows


$$\mathbf{X} = \underbrace{\begin{bmatrix} -\mathbf{x}_1^T- \\ -\mathbf{x}_2^T- \\ \vdots \\ -\mathbf{x}_N^T- \end{bmatrix}}_{\text{input data matrix}}, \quad \mathbf{y} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}}_{\text{target vector}}.$$

- 2: Compute the pseudo-inverse  $\mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$ .
- 3: Return  $\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$ .

<http://work.caltech.edu/slides/slides03.pdf>

## Boston housing dataset

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33



<https://www.cs.toronto.edu/~delve/data/boston/bostonDetail.html>

Show me the code

```
w = np.linalg.pinv(Xtr).dot(Ytr)
pred = Xte.dot(w)
loss = np.mean((pred-Yte) ** 2)
```

vectorized computation

<https://colab.research.google.com/drive/1GIovpb0ij4bSK1-jPKjNTlmXHSwTWwou>

Computational complexity?


$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$O(nd^2 + d^3)$$

Training might be computationally expensive

Nonlinear features

Data is not always 'linear'



## Transforming the data

- ▶ Linear regression => **linear in the weights**
  - ✓ linear combination of the features
- ▶ Nonlinear functions
  - ✓ can transform the data nonlinearly using any feature transformations

$$\mathbf{x} = (x_0, \dots, x_d) \xrightarrow{\Phi} \mathbf{z} = (x_0, \dots, z_{\tilde{d}})$$

input space  $\mathcal{X} = \mathbb{R}^{d+1}$       feature space  $\mathcal{Z} = \mathbb{R}^{\tilde{d}+1}$

## Transforming the data

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_{\tilde{d}}(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ z_1 \\ \vdots \\ z_{\tilde{d}} \end{bmatrix}$$

$$h(\mathbf{x}) = \tilde{\mathbf{w}}^T \Phi(\mathbf{x})$$

## Polynomial models on one feature

- ▶ A **k-th** order polynomial model in one variable is defined as:

$$h(\mathbf{x}) = w_0 + w_1 x^1 + w_2 x^2 + \dots + w_k x^k$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_k(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x^1 \\ \vdots \\ x^k \end{bmatrix}$$

## Polynomial models on two features

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \Phi_2(\mathbf{x}) \\ \Phi_3(\mathbf{x}) \\ \Phi_4(\mathbf{x}) \\ \Phi_5(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1^2 \\ x_1 x_2 \\ x_2^2 \end{bmatrix}$$



## Show me the code

```
# this function also adds the column of +1s  
poly = PolynomialFeatures(p)
```

```
# transform data
```

```
_xtr = poly.fit_transform(Xtr)
```

```
_xte = poly.fit_transform(Xte)
```

```
# linear regression
```

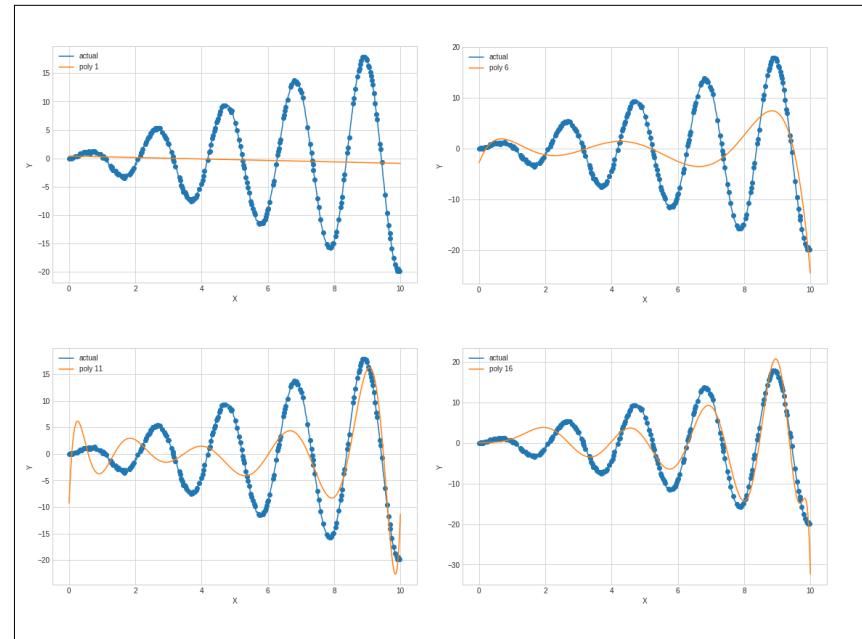
```
w = np.linalg.pinv(_xtr).dot(Ytr)
```

```
# record losses
```


```
train_loss = np.mean((_xtr.dot(w)-Ytr)**2)
```

```
test_loss = np.mean((_xte.dot(w)-Yte)**2)
```

[https://colab.research.google.com/drive/1W9kR\\_cbjYw0Ek2rsTO7\\_ojbfzxVN3pSJ#scrollTo=Wlm7SPzqhWnP](https://colab.research.google.com/drive/1W9kR_cbjYw0Ek2rsTO7_ojbfzxVN3pSJ#scrollTo=Wlm7SPzqhWnP)



## Trying a few transformations



## Polynomial models on more features

### ► **PolynomialFeatures** from *scikit-learn*

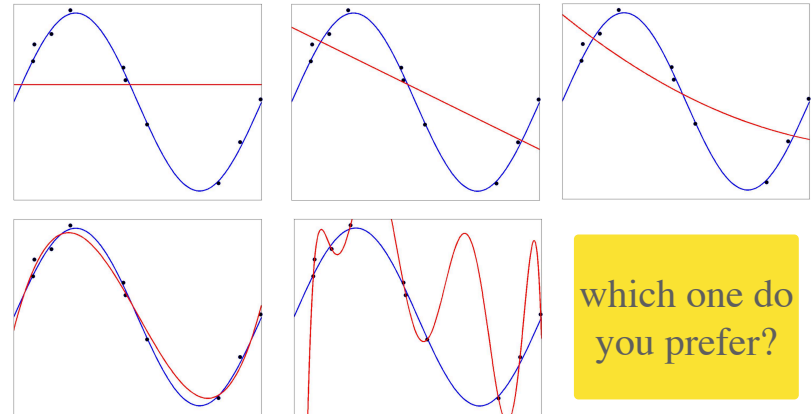
- ✓ “all polynomial combinations of the features with degree less than or equal to the specified degree”

### ► Transformation function can be anything

- ✓ choose transformation **before** looking into the data
- ✓ use **cross-validation**
- ✓ be aware of **computational cost**
- ✓ be aware of **overfitting**


# Overfitting and Regularization

## Lets talk about overfitting




Underfitting: model is too simple  
Overfitting: model is too complex

## Model complexity



## What if we restrict the hypothesis space?



## Regularization

- ▶ Adding a **penalty** to the weights to control the complexity of the model
  - ✓ usually penalizing higher weights (**except intercept**)
  - ✓ results in **simpler** or **more sparse** solutions
- ▶ Impact of regularization can be controlled by a parameter (*lambda*)

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda R(\mathbf{w})$$

## L2 regularization

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

a.k.a. Ridge Regression

Can solve using matrix calculus again:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$


always invertible

## L2 regularization

- ▶ If using the closed form solution for regularization the top-left corner of the identity matrix can be set to 0 (to handle intercept)

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix}$$

## How does it work?



## L1 regularization


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_1$$

a.k.a. Lasso Regression

- ▶ Lasso does not have a **closed form solution**
  - ▶ can solve with quadratic programming or variants of gradient descent (subgradient methods)
- ▶ The regularization term is not differentiable

## Comparison

- ▶ L1 regularization tends to generate **sparser** solutions



## Final remarks

- ▶ Linear regression
  - ✓ solved by defining a hypothesis space and a loss function
  - ✓ essentially an optimization problem that can be solved directly (closed-form) or using other techniques, such as, gradient descent
- ▶ Important concepts
  - ✓ nonlinear features
  - ✓ overfitting/underfitting
  - ✓ regularization