

Linear Regression


CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez
University of Rhode Island

Continuous targets

- Certain applications require the prediction of continuous values





<https://towardsdatascience.com/stock-prediction-in-python-b66555171a2>

Linear functions

- Assumes the output y is a **linear function** of the input x

✓ can use the function to make predictions, very simple approach, e.g. linear regression




$$y = wx + b \quad w, x, b \in \mathbb{R}$$

Draw the models


$$y = wx + b$$

can also use a vector
 $\mathbf{w} = [w_0, w_1]^T$,
where $w_0 = b$


$$y = w_1x_1 + w_0$$



$$\mathbf{w} = [0.5, 0]$$




$$\mathbf{w} = [0, 1.5]$$



$$\mathbf{w} = [0.5, 0.5]$$

Linear functions

► What if we have d features?



why transpose?

$$y = \mathbf{w}^T \mathbf{x} + b \quad \mathbf{w}, \mathbf{x} \in \mathbb{R}^d \quad b \in \mathbb{R}$$

Figure from <https://www.statlearning.com/>

Linear functions

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

hypothesis

weights

bias

The weights and bias are the model **parameters** which define the hypothesis and are used to make predictions

Alternative notation

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = \sum_{i=1}^d w_i x_i + b$$

x0	x1	x2	y
1	0.5	0.1	0.25
1	0.3	0.9	0.5
1	0.3	0.875	1.15
1	0.45	0.15	2.13
...

the bias can be absorbed into the summation if we augment \mathbf{w} and \mathbf{x}

$$\begin{aligned} h(\mathbf{x}) &= \sum_{i=0}^d \tilde{w}_i \tilde{x}_i \\ &= \tilde{\mathbf{w}}^T \tilde{\mathbf{x}} \end{aligned}$$

Augmented vectors

input: $(\overset{+1}{x_0}, x_1, \dots, x_d)$

model: (w_0, w_1, \dots, w_d)


$$h(\mathbf{x}) = \sum_{i=0}^d w_i x_i = \mathbf{w}^T \mathbf{x}$$

What is the hypothesis space?

all linear functions in \mathbb{R}^{d+1}

$$\mathcal{H} = \{h_{\mathbf{w}} : h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x}, \mathbf{w} \in \mathbb{R}^{d+1}\}$$


Residuals



Want a **linear function** with **small residuals**

<http://work.caltech.edu/slides/slides03.pdf>

Residuals and Loss



High loss

Low loss

From <https://developers.google.com/machine-learning/crash-course/descending-into-ml/training-and-loss>

Goal of learning

- Find a function that best approximates target function (**minimize expected loss**)

For $h \in \mathcal{H}$ and $\forall (\mathbf{x}^{(i)}, y^{(i)}) \sim P$, we want $h(\mathbf{x}^{(i)}) \approx f(\mathbf{x}^{(i)})$

- What is the **expected loss**?

✓ cannot calculate, can **approximate** with empirical loss:

$$\mathbb{E} [l(h, \mathbf{x}^{(i)}, y^{(i)})]_{(\mathbf{x}^{(i)}, y^{(i)}) \sim P} \approx L(h, \mathcal{D})$$

Defining linear regression

- Data $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$

$$\mathbf{x}^{(i)} \in \mathbb{R}^{d+1}, y^{(i)} \in \mathbb{R}$$

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$$

- Loss Function: Squared Loss (MSE)


$$l_{sq}(h, \mathbf{x}, y) = (h(\mathbf{x}) - y)^2$$

$$L_{sq}(h, \mathcal{D}) = \frac{1}{n} \sum_{i=1}^n l_{sq}(h, \mathbf{x}^{(i)}, y^{(i)})$$

Closed-form Solution

Normal Equation

min vs. argmin



$\min f(x)?$

$\arg \min_x f(x)?$

$$f(x) = 6x^2 + 4x + 5$$

Solving linear regression (least squares)

find these parameters

$\mathbf{w}^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$
given this objective function

unconstrained optimization

Solving linear regression (least squares)

$$\mathbf{w}^* = \arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (h(\mathbf{x}^{(i)}) - y^{(i)})^2$$

↓ $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$

Understanding matrix form

$$\mathbf{y} = \mathbf{X}\mathbf{w}$$

$$\begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix} \approx \begin{bmatrix} (\mathbf{x}^{(1)})^T \\ \vdots \\ (\mathbf{x}^{(n)})^T \end{bmatrix} \cdot \begin{bmatrix} w_0 \\ \vdots \\ w_d \end{bmatrix}$$

input vector $\mathbf{x}^{(1)}$

Norms

- ▶ A **norm** is a function that assigns a strictly positive length to each vector in a vector space
✓ except for the zero vector

ℓ_1 -norm:
manhattan distance
from origin

$$\|\mathbf{x}\|_1 = \sum_{i=1}^n |x_i|$$

ℓ_2 -norm:
euclidean norm, euclidean
distance from origin

$$\|\mathbf{x}\|_2 = \left(\sum_{i=1}^n x_i^2 \right)^{\frac{1}{2}}$$

Using matrix notation


$$\arg \min_{h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (\mathbf{w}^T \mathbf{x}^{(i)} - y^{(i)})^2$$



$$\arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$

How to minimize it?


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2$$



continuous,
differentiable,
convex

**optimal
solution**

(some) Critical points



https://en.wikipedia.org/wiki/Maxima_and_minima

Closed form solution

$$E_{\text{in}}(\mathbf{w}) = \frac{1}{N} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E_{\text{in}}(\mathbf{w}) = \frac{2}{N} \mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = \mathbf{0}$$

$$\mathbf{X}^T \mathbf{X} \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} \quad \text{where} \quad \mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$

\mathbf{X}^\dagger is the 'pseudo-inverse' of \mathbf{X}

There are other methods for finding the optimal solution
e.g. gradient descent, MLE

<http://work.caltech.edu/slides/slides03.pdf>

The algorithm

1. Construct the matrix \mathbf{X} and the vector \mathbf{y} from the data set $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ as follows

$$\mathbf{X} = \underbrace{\begin{bmatrix} -\mathbf{x}_1^T \\ -\mathbf{x}_2^T \\ \vdots \\ -\mathbf{x}_N^T \end{bmatrix}}_{\text{input data matrix}}, \quad \mathbf{y} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}}_{\text{target vector}}.$$

2. Compute the pseudo-inverse $\mathbf{X}^\dagger = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$.
3. Return $\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$.

<http://work.caltech.edu/slides/slides03.pdf>

Show me the code

```
w = np.linalg.pinv(Xtr).dot(Ytr)
```

```
pred = Xte.dot(w)
```

```
loss = np.mean((pred-Yte) ** 2)
```

vectorized computation

Computational complexity?

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$nd^2 + d^3 + nd + d^2$$

$$O(nd^2 + d^3)$$

Training might be computationally expensive

Colab notebook

<https://colab.research.google.com/drive/1GIovpb0ij4bSK1-jPKjNTlmXHSwTWwou>