

Nonlinear features, Regularization


CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez
University of Rhode Island

Nonlinear features

Data is not always 'linear'



Transforming the data

▸ Linear regression => **linear in the weights**

✓ linear combination of the features

▸ Nonlinear functions

✓ can transform the data nonlinearly using any feature transformations

$$\begin{array}{ccc} \mathbf{x} = (x_0, \dots, x_d) & \xrightarrow{\Phi} & \mathbf{z} = (x_0, \dots, z_{\tilde{d}}) \\ \text{input space } \mathcal{X} = \mathbb{R}^{d+1} & & \text{feature space } \mathcal{Z} = \mathbb{R}^{\tilde{d}+1} \end{array}$$

Transforming the data

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_{\tilde{d}}(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ z_1 \\ \vdots \\ z_{\tilde{d}} \end{bmatrix}$$

$$h(\mathbf{x}) = \tilde{\mathbf{w}}^T \Phi(\mathbf{x})$$

Polynomial models on one feature

► A **k-th** order polynomial model in one variable is defined as:

$$h(\mathbf{x}) = w_0 + w_1x^1 + w_2x^2 + \dots + w_kx^k$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_k(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x^1 \\ \vdots \\ x^k \end{bmatrix}$$

Polynomial models on two features

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \Phi_2(\mathbf{x}) \\ \Phi_3(\mathbf{x}) \\ \Phi_4(\mathbf{x}) \\ \Phi_5(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1^2 \\ x_1x_2 \\ x_2^2 \end{bmatrix}$$


Show me the code

```
# this function also adds the column of +1s
poly = PolynomialFeatures(p)


# transform data
_xtr = poly.fit_transform(Xtr)
_xte = poly.fit_transform(Xte)

# linear regression
w = np.linalg.pinv(_xtr).dot(Ytr)

# record losses
train_loss = np.mean((_xtr.dot(w)-Ytr)**2)
test_loss = np.mean((_xte.dot(w)-Yte)**2)
```



Trying a few transformations

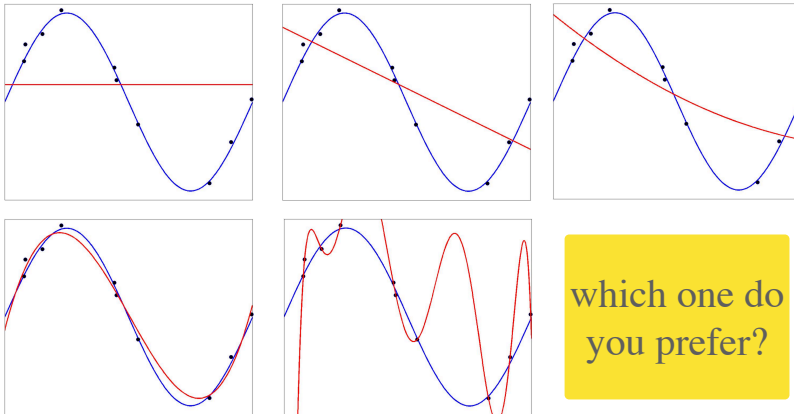


Polynomial models on more features

- ▶ **PolynomialFeatures** from *scikit-learn*
 - ✓ “all polynomial combinations of the features with degree less than or equal to the specified degree”
- ▶ Transformation function can be anything
 - ✓ choose transformation **before** looking into the data
 - ✓ use **cross-validation**
 - ✓ be aware of **computational cost**
 - ✓ be aware of **overfitting**

Overfitting and
Regularization


Lets talk about overfitting




which one do you prefer?

Underfitting: model is too simple
Overfitting: model is too complex

Model complexity



What if we restrict the hypothesis space?



Regularization

- Adding a **penalty** to the weights to control the complexity of the model
 - ✓ usually penalizing higher weights (**except intercept**)
 - ✓ results in **simpler** or **more sparse** solutions
- Impact of regularization can be controlled by a parameter (*lambda*)

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda R(\mathbf{w})$$

L2 regularization

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

a.k.a. Ridge Regression

Can solve using matrix calculus again:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$


always invertible

L2 regularization

- If using the closed form solution for regularization the top-left corner of the identity matrix can be set to 0 (to handle intercept)

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix}$$

How does it work?



<http://work.caltech.edu/slides/slides12.pdf>

L1 regularization


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_1$$

a.k.a. Lasso Regression

- Lasso does not have a **closed form solution**
 - can solve with quadratic programming or variants of gradient descent (subgradient methods)
- The regularization term is not differentiable

Comparison

- ▶ L1 regularization tends to generate **sparser** solutions



Final remarks

- ▶ Linear regression
 - ✓ solved by defining a hypothesis space and a loss function
 - ✓ essentially an optimization problem that can be solved directly (closed-form) or using other techniques, such as, gradient descent
- ▶ Important concepts
 - ✓ nonlinear features
 - ✓ overfitting/underfitting
 - ✓ regularization

Colab notebook

https://colab.research.google.com/drive/1W9kR_cbjYw0Ek2rsTO7_ojbfzxVN3pSJ#scrollTo=Wlm7SPzqhWnP