

# Nonlinear features, Regularization

CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez  
University of Rhode Island

## Linear regression (closed form solution)

- 1: Construct the matrix  $\mathbf{X}$  and the vector  $\mathbf{y}$  from the data set  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$  as follows

$$\mathbf{X} = \underbrace{\begin{bmatrix} -\mathbf{x}_1^\top - \\ -\mathbf{x}_2^\top - \\ \vdots \\ -\mathbf{x}_N^\top - \end{bmatrix}}_{\text{input data matrix}}, \quad \mathbf{y} = \underbrace{\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}}_{\text{target vector}}.$$

- 2: Compute the pseudo-inverse  $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$ .
- 3: Return  $\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$ .


<http://work.caltech.edu/slides/slides03.pdf>

## A note on the pseudoinverse

- ▶ If the inverse  $A^{-1}$  of a matrix exists ...
  - ✓  $AA^{-1} = A^{-1}A = I$
  - ✓ square and full-rank
- ▶ ... then the pseudoinverse is the inverse
  - ✓  $A^\dagger = A^{-1}$

# Nonlinear features

## Data is not always 'linear'



## Transforming the data

- ▶ Linear regression => **linear in the weights**
  - ✓ linear combination of the features
- ▶ Nonlinear functions
  - ✓ can transform the data nonlinearly using any feature transformations

$$\mathbf{x} = (x_0, \dots, x_d) \xrightarrow{\Phi} \mathbf{z} = (x_0, \dots, z_{\tilde{d}})$$

input space  $\mathcal{X} = \mathbb{R}^{d+1}$       feature space  $\mathcal{Z} = \mathbb{R}^{\tilde{d}+1}$

## Transforming the data

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_{\tilde{d}}(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ z_1 \\ \vdots \\ z_{\tilde{d}} \end{bmatrix}$$

$$h(\mathbf{x}) = \tilde{\mathbf{w}}^T \Phi(\mathbf{x})$$

## Polynomial features

- ▶ A **k-th** order polynomial transformation on **one** variable:

$$\mathbf{x} = \begin{bmatrix} 1 \\ x \end{bmatrix}$$

$$\Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_k(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x^1 \\ \vdots \\ x^k \end{bmatrix}$$

## Polynomial models on one feature

- ▶ A **k-th** order polynomial model on one variable can be defined as:

$$h(\mathbf{x}) = w_0 + w_1x^1 + w_2x^2 + \dots + w_kx^k$$

$$\mathbf{x} = \begin{bmatrix} 1 \\ x \end{bmatrix} \quad \Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \vdots \\ \Phi_k(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x^1 \\ \vdots \\ x^k \end{bmatrix}$$

## Polynomial features on two variables

$$\mathbf{x} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}$$

$$\Phi(\mathbf{x}) = \mathbf{z} = \begin{bmatrix} 1 \\ \Phi_1(\mathbf{x}) \\ \Phi_2(\mathbf{x}) \\ \Phi_3(\mathbf{x}) \\ \Phi_4(\mathbf{x}) \\ \Phi_5(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \\ x_1^2 \\ x_1x_2 \\ x_2^2 \end{bmatrix}$$

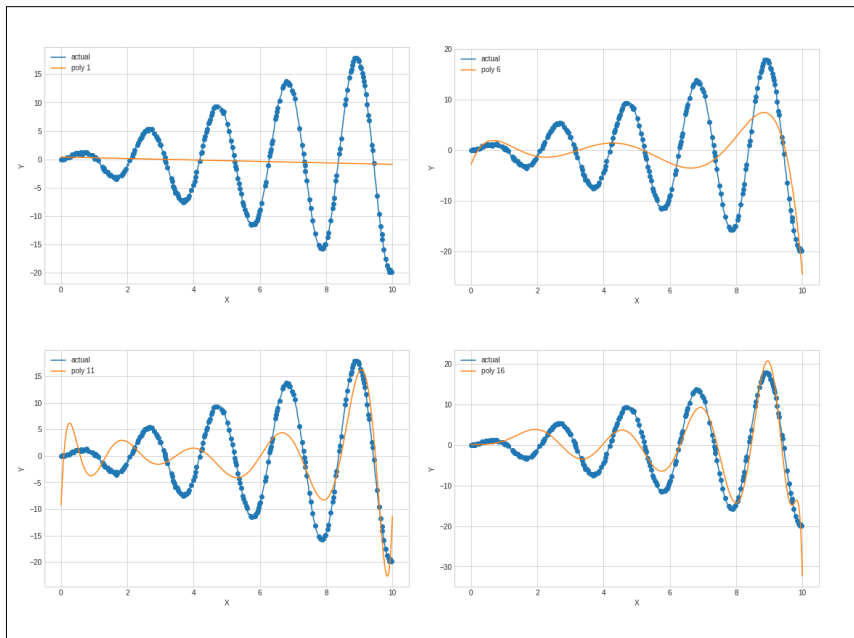
## Show me the code

```
# this function also adds the column of +1s  
poly = PolynomialFeatures(p)
```


```
# transform data  
_xtr = poly.fit_transform(Xtr)  
_xte = poly.fit_transform(Xte)
```

```
# linear regression  
w = np.linalg.pinv(_xtr).dot(Ytr)
```

```
# record losses  
train_loss = np.mean((_xtr.dot(w) - Ytr)**2)  
test_loss = np.mean((_xte.dot(w) - Yte)**2)
```



## Trying a few transformations

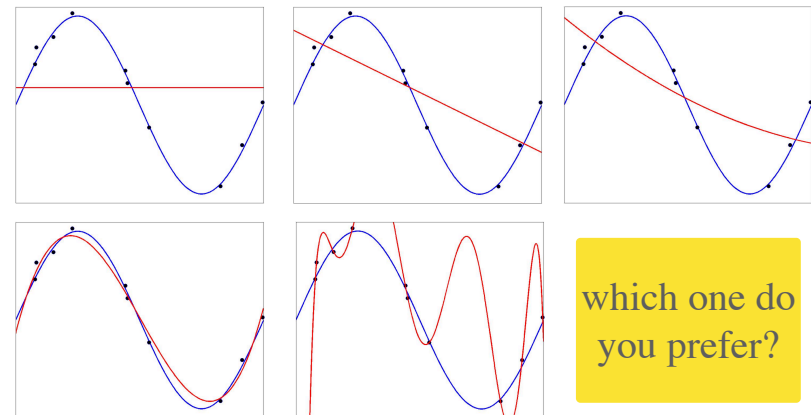


## Feature transformations

- ▶ **PolynomialFeatures** from *scikit-learn*
  - ✓ “all polynomial combinations of the features with degree less than or equal to the specified degree”
- ▶ Transformation function can be anything
  - ✓ choose transformation **before** looking into the data
  - ✓ use **cross-validation**
  - ✓ be aware of **computational cost**
  - ✓ be aware of **overfitting**


## Overfitting and Regularization

## Overfitting and hypothesis spaces



Underfitting: model is too simple  
Overfitting: model is too complex


## Model complexity



## Regularization

- ▶ For polynomial features/models
  - ✓ the degree **k** of the polynomial controls the complexity of the model
  - ✓ usually a hyperparameter search is necessary for finding the right complexity
- ▶ Alternative approach
  - ✓ keep the hypothesis space larger, but **regularize** it

## What if we restrict the hypothesis space?



## Regularization

- ▶ Adding a **penalty** to the weights to control the complexity of the model
  - ✓ usually penalizing higher weights (**except intercept**)
  - ✓ results in **simpler** or **more sparse** solutions
- ▶ Impact of regularization can be controlled by a parameter (*lambda*)

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda R(\mathbf{w})$$

## L2 regularization

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

a.k.a. Ridge Regression

Can solve using matrix calculus again:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$


always invertible

## L2 regularization

- If using the closed form solution for regularization the **top-left** corner of the identity matrix can be set to 0 (to handle intercept)

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{bmatrix}$$

## How does it work?



<http://work.caltech.edu/slides/slides12.pdf>

## L1 regularization


$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{n} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_1$$

a.k.a. Lasso Regression

- Lasso does not have a **closed form solution**
  - can solve with quadratic programming or variants of gradient descent (subgradient methods)
- The regularization term is not differentiable

## Comparison

- ▶ L1 regularization tends to generate **sparser** solutions



## Final remarks

- ▶ Linear regression
  - ✓ solved by defining a hypothesis space and a loss function
  - ✓ essentially an optimization problem that can be solved directly (closed-form) or using other techniques, such as, gradient descent
- ▶ Important concepts
  - ✓ nonlinear features
  - ✓ overfitting/underfitting
  - ✓ regularization

## Colab notebook

[https://colab.research.google.com/drive/1W9kR\\_cbjYw0Ek2rsTO7\\_ojbfzxVN3pSJ#scrollTo=Wlm7SPzqhWnP](https://colab.research.google.com/drive/1W9kR_cbjYw0Ek2rsTO7_ojbfzxVN3pSJ#scrollTo=Wlm7SPzqhWnP)