

Linear Classifiers, Logistic Regression

CSC 461: Machine Learning

Fall 2022

Prof. Marco Alvarez
University of Rhode Island

Linear classifiers

Linear classifiers

- ▶ Discriminative

- ✓ Perceptron

- ✓ Logistic regression

- ✓ Support vector machines

- ▶ Generative

- ✓ Linear discriminant analysis

- ✓ Naive bayes

Binary classification



$$\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^d$$

$$y^{(i)} \in \{-1, +1\}$$


x1	x2	y
0.5	0.1	+1
0.3	0.9	-1
0.3	0.875	-1
0.45	0.15	+1
...

Plots (regression x classification)




Binary classification goal

- ▶ Learn a **decision boundary** such that two classes can be separated




The *sign* function

$$\text{sign}(x) = \begin{cases} -1 & \text{if } x < 0, \\ 0 & \text{if } x = 0, \\ +1 & \text{if } x > 0 \end{cases}$$




$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$$

Decision boundary



Decision boundary

A hyperplane in \mathbb{R}^2 is a line



$$0 = b + w_1x_1 + w_2x_2$$

$$x_2 = -\frac{b}{w_2} - \frac{w_1}{w_2}x_1$$

Absorbing the bias

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$$

$$= \text{sign}\left(\sum_{i=1}^d w_i x_i + b\right)$$


$$x_0 = 1, \quad w_0 = b$$

$$h(\mathbf{x}) = \text{sign}\left(\sum_{i=0}^d w_i x_i\right)$$


$$= \text{sign}(\mathbf{w}^T \mathbf{x})$$

x0	x1	x2	Y
1	0.5	0.1	+1
1	0.3	0.9	-1
1	0.3	0.875	-1
1	0.25	0.561	-1
1	0.45	0.15	+1
...

Absorbing the bias



Learning




Example

- Provide a solution (weight vector)

x_0	x_1	x_2	y
1	0	0	-1
1	0	1	-1
1	1	0	-1
1	1	1	+1

The *sign* function (again)


$$\text{sign}(x) = \begin{cases} -1 & \text{if } x < 0, \\ 0 & \text{if } x = 0, \\ +1 & \text{if } x > 0 \end{cases}$$



Note that the gradient is zero almost everywhere and the gradient is undefined at $x = 0$.

Can we use the squared loss?

- Treat target labels (binary) as continuous
 - ✓ final prediction decided by checking $h(\mathbf{x}) > 0$




Predicted values can fall outside $[-1, 1]$ range

Square loss penalizes correct predictions with large losses

Logistic regression

Logistic function

$$\sigma(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$



mapping \mathbb{R} to $[0,1]$

continuous and
differentiable

Logistic regression

- ▶ Binary classifier

- ✓ uses a **logistic function** (type of sigmoid function, S-shaped)
- ✓ models **probability** of output in terms of input


- ▶ It is considered a **linear classifier**

- ✓ even though the *activation function* is non-linear

- ▶ It is a **discriminative model**

- ✓ models decision boundary directly, $P(y | \mathbf{x})$ in this case

Conditional probabilities



$$P(a|b) = \frac{P(a,b)}{P(b)}$$

$P(X, Y)$

X	Y	P
+x	+y	0.2
+x	-y	0.3
-x	+y	0.4
-x	-y	0.1

$P(+x | +y) ? .2 / .6$ $P(-x | +y) ? .4 / .6$ $P(-y | +x) ? .3 / .5$

Set up

$$\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$$

$$\mathbf{x}^{(i)} \in \mathbb{R}^d$$

$$y^{(i)} \in \{-1, +1\}$$

Probabilistic interpretation

$$h(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} = \frac{e^{\mathbf{w}^T \mathbf{x}}}{e^{\mathbf{w}^T \mathbf{x}} + 1}$$

(probability of class +1)

$$P(y = +1 \mid \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} = \sigma(\mathbf{w}^T \mathbf{x})$$

$$P(y = -1 \mid \mathbf{x}) = 1 - P(y = +1 \mid \mathbf{x})$$

(probability of class -1)

$$P(y = -1 \mid \mathbf{x}) = \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}}} = \sigma(-\mathbf{w}^T \mathbf{x})$$

Probabilistic interpretation

(probability of class +1)

$$P(y = +1 \mid \mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}}} = \sigma(\mathbf{w}^T \mathbf{x})$$

(probability of class -1)


$$P(y = -1 \mid \mathbf{x}) = \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}}} = \sigma(-\mathbf{w}^T \mathbf{x})$$

$$P(y \mid \mathbf{x}) = \frac{1}{1 + e^{-y\mathbf{w}^T \mathbf{x}}} = \sigma(y\mathbf{w}^T \mathbf{x})$$

Decision boundary


$$P(y = +1 \mid \mathbf{x}) = P(y = -1 \mid \mathbf{x}) = 0.5$$

$\sigma(\mathbf{w}^T \mathbf{x})$ $\sigma(-\mathbf{w}^T \mathbf{x})$

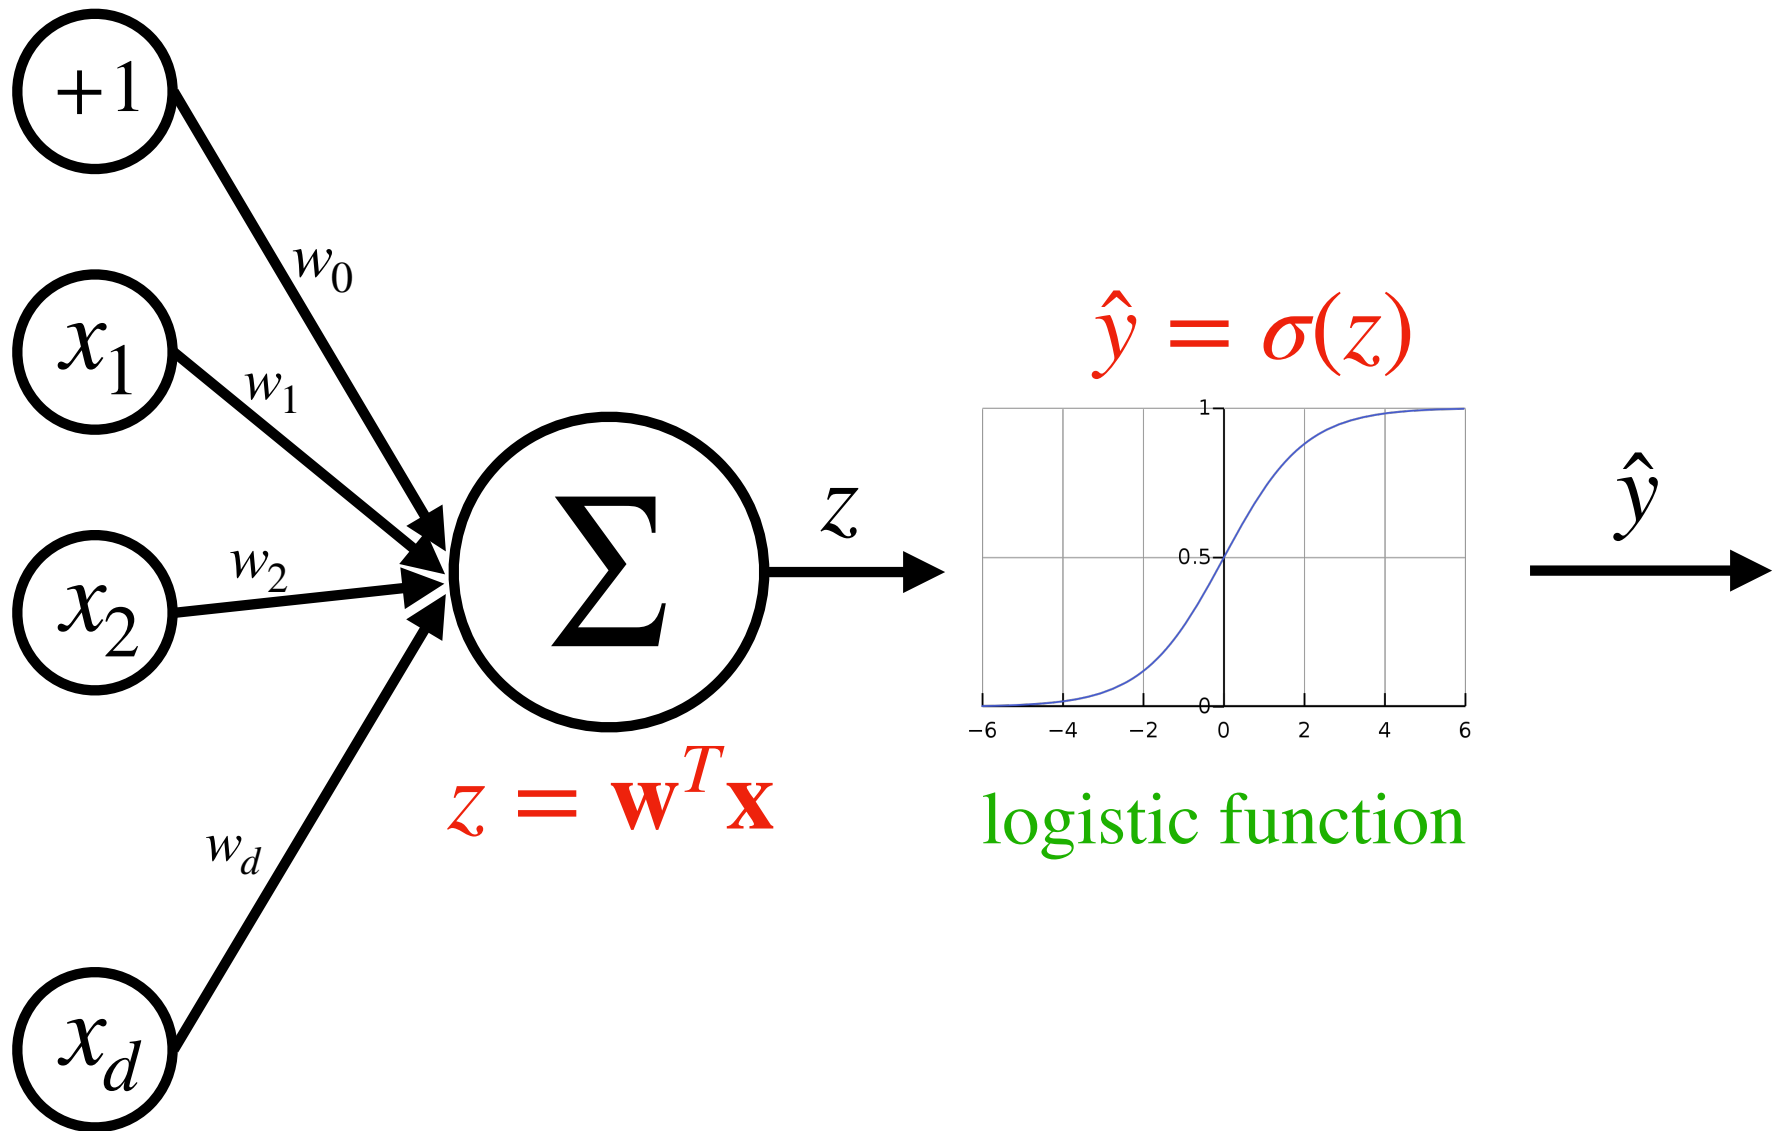


Logistic regression has a linear decision boundary $\mathbf{w}^T \mathbf{x} = 0$

Linear decision boundary



Logistic regression



Solving logistic regression

Maximum likelihood estimation (MLE)

- ▶ MLE estimates **parameters** based on the principle that if we observe \mathcal{D} , we should choose the parameters that make \mathcal{D} most probable
- ▶ We can derive formulas for W that maximize $p(\mathcal{D}; W)$
 - ✓ many machine learning algorithms follow this **maximum likelihood** principle
 - ✓ **want** $P(y \mid \mathbf{x}; W)$
 - ✓ **learn** $W^* = \arg \max_{\mathbf{W}} P(y \mid \mathbf{x}; W)$

Maximum likelihood estimation (MLE)

- ▶ The **conditional data likelihood** $\mathcal{L}(\mathbf{w})$ is the probability of the observed labels y conditioned on the feature values \mathbf{x}

✓ weights can be learned by maximizing this likelihood

$$\mathcal{L}(\mathbf{w}) = P(y^{(1)}, \dots, y^{(n)} \mid \mathbf{x}^{(1)}, \dots, \mathbf{x}^{(n)}; \mathbf{w}) = \prod_{i=1}^n P(y^{(i)} \mid \mathbf{x}^{(i)}; \mathbf{w})$$

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \prod_{i=1}^n P(y^{(i)} \mid \mathbf{x}^{(i)}; \mathbf{w})$$

Maximum likelihood estimation

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \prod_{i=1}^n P(y^{(i)} \mid \mathbf{x}^{(i)}; \mathbf{w})$$

$$= \arg \max_{\mathbf{w}} \log \left(\prod_{i=1}^n P(y^{(i)} \mid \mathbf{x}^{(i)}; \mathbf{w}) \right)$$

$$= \arg \max_{\mathbf{w}} \frac{1}{n} \log \left(\prod_{i=1}^n P(y^{(i)} \mid \mathbf{x}^{(i)}; \mathbf{w}) \right) \quad \frac{1}{1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}}}$$

$$= \arg \max_{\mathbf{w}} -\frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right)$$

negative log
likelihood

$$= \arg \min_{\mathbf{w}} \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right)$$

error (loss)
 $e(h(\mathbf{x}), y)$


Loss function

$$L(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right)$$


cross-entropy loss
(over a dataset)

- ▶ no closed-form solution (non-linear function), but loss is convex
- ▶ can use gradient descent or second-order methods

Example: 2d dataset




Example: loss function



plot shows contour
lines in the space
of parameters w_1
and w_2 ,
 w_0 is omitted

Solution



Logistic function (derivative)

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma'(x) = \frac{(1 + e^{-x})(0) - (1)(-e^{-x})}{(1 + e^{-x})^2}$$

$$= \frac{e^{-x}}{(1 + e^{-x})^2}$$

$$= \sigma(x)(1 - \sigma(x))$$

Gradient

$$\nabla_{\mathbf{w}} L(\mathbf{w}) = \left[\frac{\partial L(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial L(\mathbf{w})}{\partial w_d} \right]$$

$$\frac{\partial L(\mathbf{w})}{\partial w_j} = \frac{\partial}{\partial w_j} \frac{1}{n} \sum_{i=1}^n \log \left(1 + e^{-y^{(i)} \mathbf{w}^T \mathbf{x}^{(i)}} \right)$$

$$= -\frac{1}{n} \sum_{i=1}^n \sigma \left(-\mathbf{w}^T \mathbf{x}^{(i)} \right) y^{(i)} x_j^{(i)}$$

How to classify new data?

- ▶ Once the final hypothesis $h(\mathbf{x})$ is known ...

- ✓ $h(\mathbf{x}) = p(+1 \mid \mathbf{x})$

- ✓ predict label $+1$ to input instance \mathbf{x}

- ✓ if $p(+1 \mid \mathbf{x}) \geq 0.5$

- ✓ predict label -1 to input instance \mathbf{x}

- ✓ if $p(+1 \mid \mathbf{x}) < 0.5$

Final remarks

- ▶ Simple classifier with **probabilistic outputs**
- ▶ Loss function is convex and can be trained with GD methods (**no closed-form**)
- ▶ Robust to overfitting
- ▶ Offers interpretability to weights (feature importance)
- ▶ However, **decision boundary is still linear**